

# Computer Architecture

## Lecture 2

Peter Mikulecky


# Integrated Approach

- Computer architecture should be viewed through an integrated approach, integrated view
- What really matters is the functioning of the complete system
  - hardware, runtime system, compiler, operating system, and application
  - In networking, this is called the “End to End argument”
- Computer architecture is not just about transistors, individual instructions, or particular implementations

# What's Computer Architecture?

The attributes of a [computing] system as seen by the programmer, i.e., the conceptual structure and functional behavior, as distinct from the organization of the data flows and controls the logic design, and the physical implementation.

Amdahl, Blaaw, and Brooks, 1964




# Trends

Gordon Moore (Founder of Intel) observed in 1965 that the number of transistors that could be crammed on a chip doubles every year.

This has CONTINUED to be true since then.

## Transistors Per Chip




# Moore's Law

- Moore's Law, first hypothesised in 1965 by Intel founder Gordon Moore, states that the number of transistors in a dense integrated circuit will double approximately every two years. The shrinking of transistors enables a larger number to be held within the same area, which results in a faster processor that can operate at lower power requirements.
- Although the law was adhered to rigidly for half a century, in 2015 Intel admitted that the pace of advancement had started to slow down. Its eighth-generation Core CPUs, codenamed Coffee Lake, are set to launch in the second half of 2017 and will once again be built on the same 14-nanometre (nm) process used three generations prior for its Broadwell chips, originally released in 2014.

# The end of the Moore's Law?


- “I don't think anyone could confidently tell you that they have a plan for 15 more years of Moore's law,” says Greg Yeric, director of future silicon technology for ARM Research.
- „Moore's Law 2017: An Uphill Battle“
- <https://eandt.theiet.org/content/articles/2017/05/moore-s-law-2017-an-uphill-battle/>
- „Is Moore's Law still the law?“
- <https://www.electronicsweekly.com/news/moores-law-still-law-2017-09/>

# Microprocessor Transistor Counts 1971-2011 & Moore's Law



# Trends

Memory Capacity (and Cost) have changed dramatically in the last 20 years.



year	size(Mb)	cyc time
1980	0.0625	250 ns
1983	0.25	220 ns
1986	1	190 ns
1989	4	165 ns
1992	16	145 ns
1996	64	120 ns
2000	256	100 ns



# Trends


Based on SPEED, the CPU has increased dramatically, but memory and disk have increased only a little. This has led to dramatic changes in architecture, operating systems, and programming practices.

	<u>Capacity</u>	<u>Speed (latency)</u>
<b>Logic</b>	<b>2x in 3 years</b>	<b>2x in 3 years</b>
<b>DRAM</b>	<b>4x in 3 years</b>	<b>2x in 10 years</b>
<b>Disk</b>	<b>4x in 3 years</b>	<b>2x in 10 years</b>

Workstation performance (measured in Spec Marks) improves roughly 50% per year (2X every 18 months)


Improvement in cost performance estimated at 70% per year

# Moore's Law 2017




<https://eandt.theiet.org/content/articles/2017/05/moore-s-law-2017-an-uphill-battle/>

# What is the next wave?



# Instruction Set Architecture: Critical Interface



- Properties of a good abstraction
  - Lasts through many generations (portability)
  - Used in many different ways (generality)
  - Provides **convenient** functionality to higher levels
  - Permits an **efficient** implementation at lower levels

# Software Abstraction

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

C

`_sum:`

```
    pushl %ebp
    movl %esp,%ebp
    movl 12(%ebp),%eax
    addl 8(%ebp),%eax
    movl %ebp,%esp
    popl %ebp
    ret
```

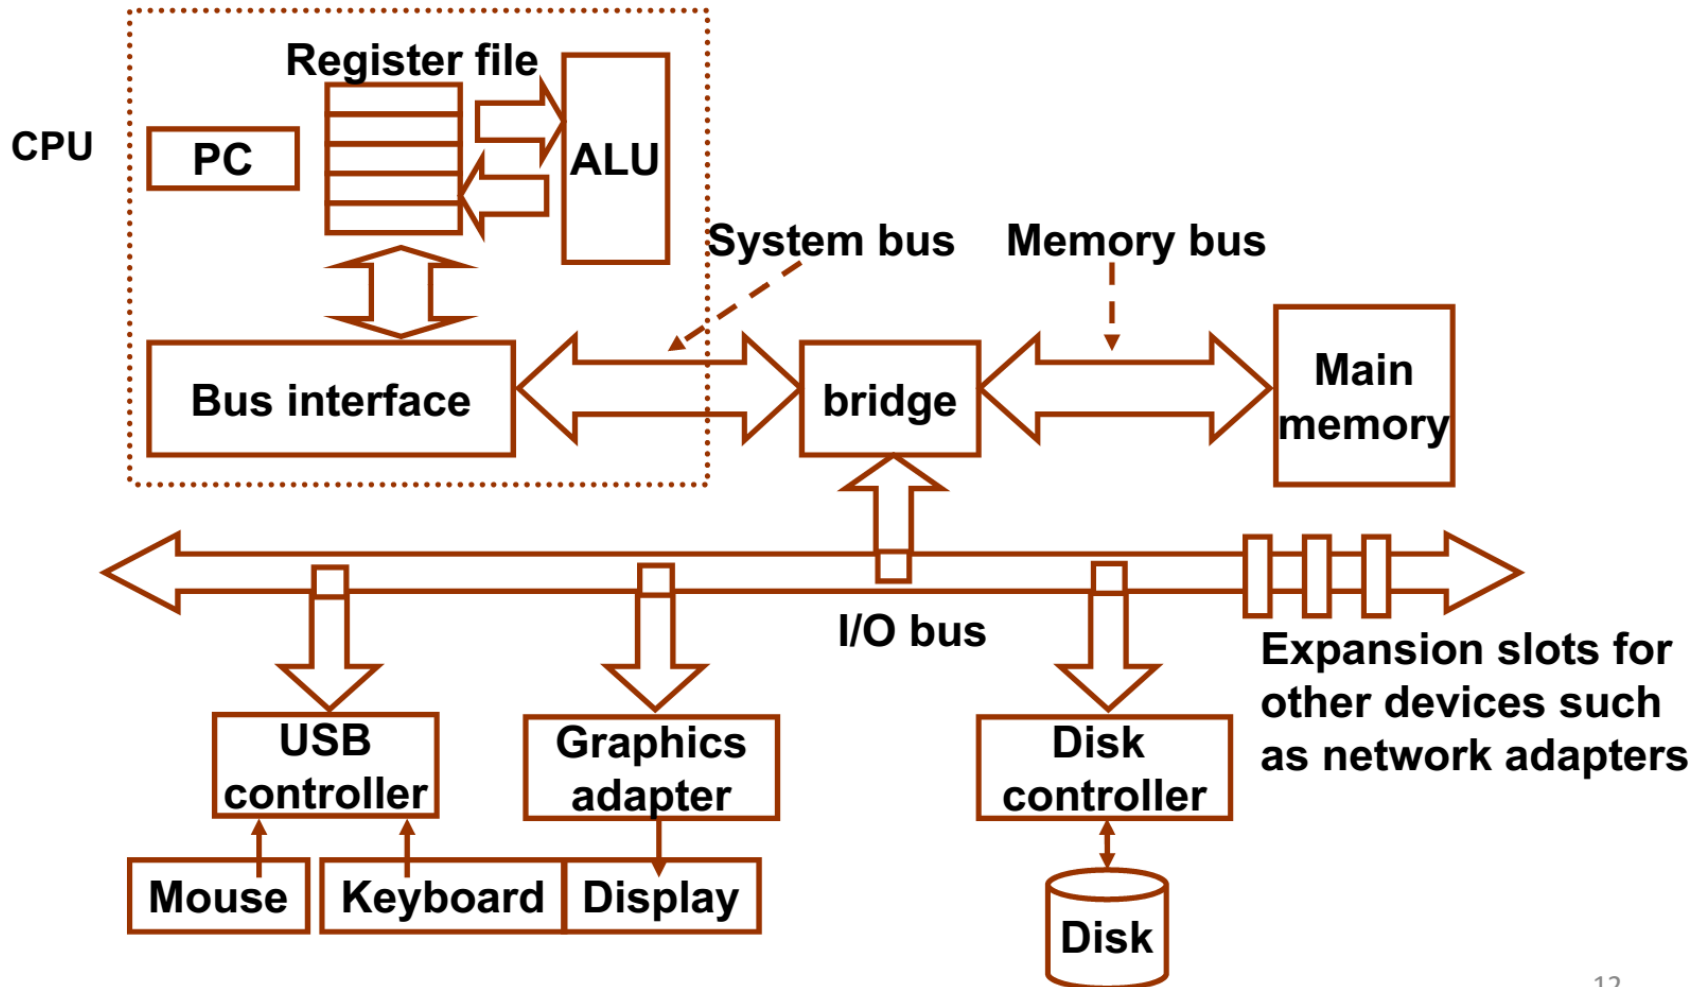
assembly

0x401040 <sum>:

0x55  
0x89  
0xe5  
0x8b  
0x45  
0x0c  
0x03  
0x45  
0x08  
0x89  
0xec  
0x5d  
0xc3


machine  
code

# Hardware Abstraction






# Real View






# Another Real View






## View 3 – Massive Cooler






# Expanding Cards





# Memories





# Hard Discs

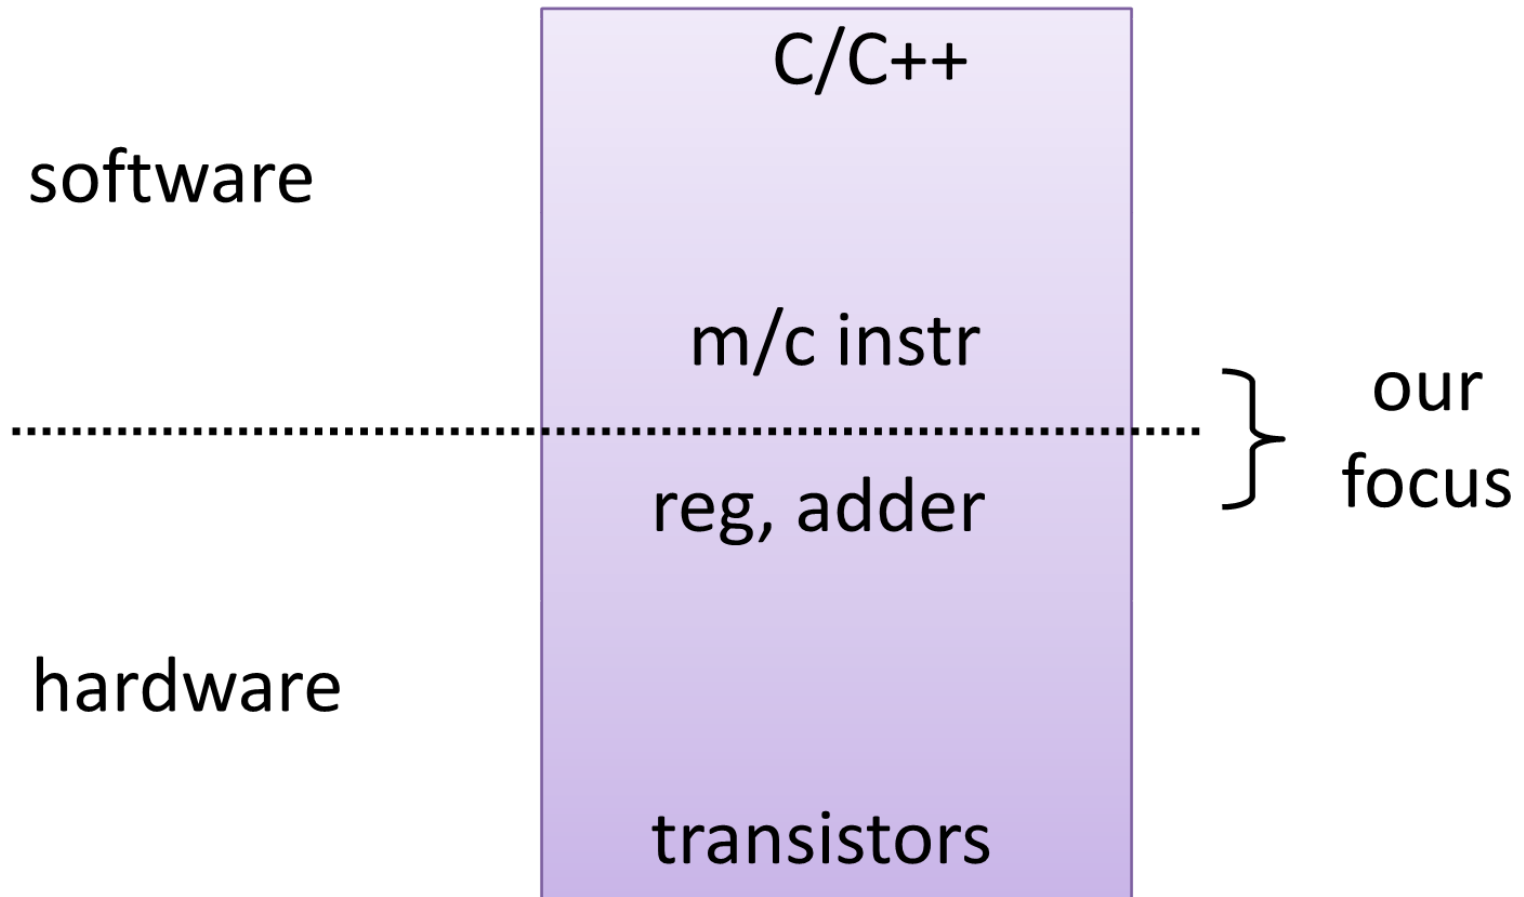




# Two Sides of a HDD



# Hardware/Software Interface

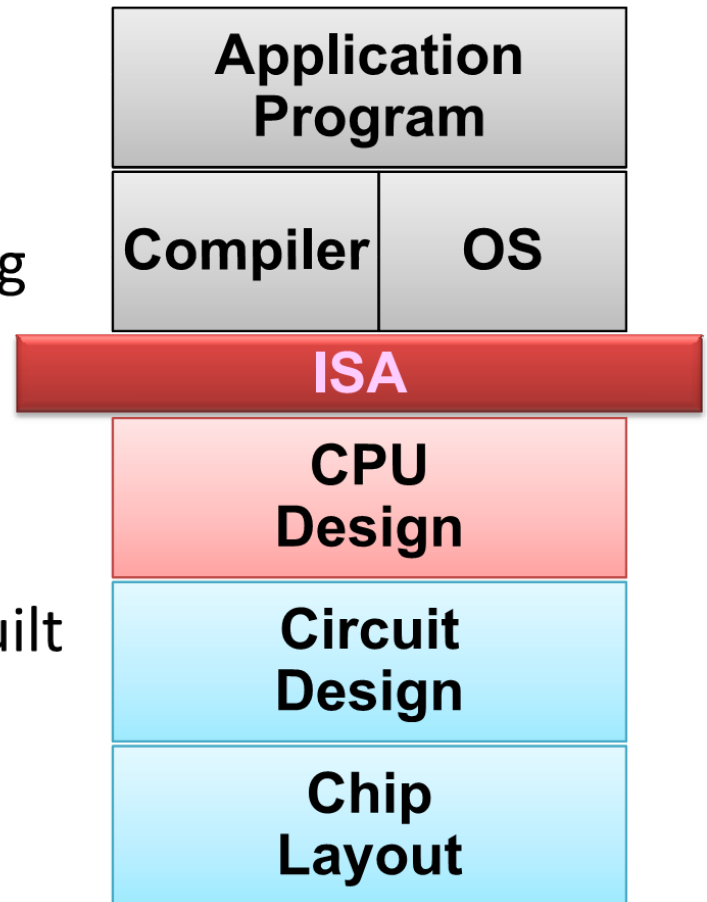


# Levels of Architecture

- Instruction set architecture (ISA)
  - Lowest level visible to a programmer
  - Operation (add/sub/mul/shift)
- Microarchitecture
  - Fills the gap between instructions and logic modules
  - Operation Vs Micro Operation

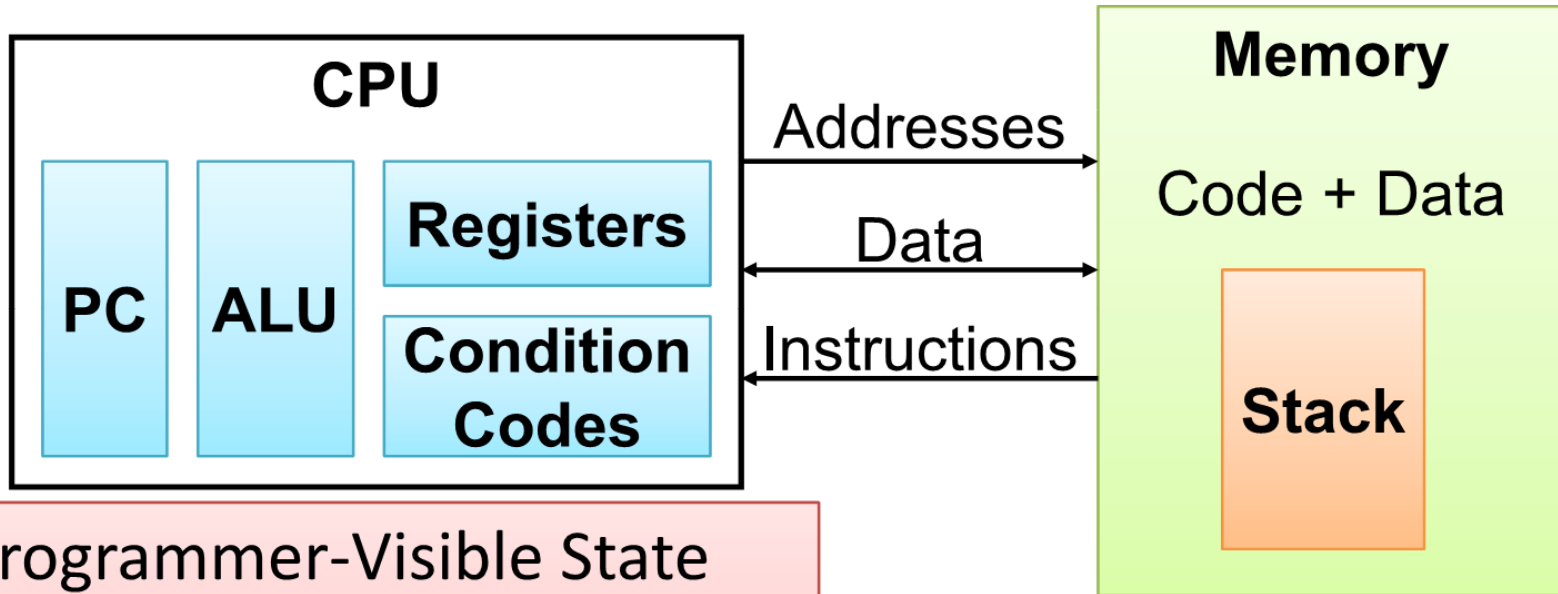
# Instruction Set Architecture

- Assembly Language View
  - Processor state (RF, mem)
  - Instruction set and encoding
- Layer of Abstraction
  - Above: how to program machine - HLL, OS
  - Below: what needs to be built
    - tricks to make it run fast





# The Abstract Machine



- Programmer-Visible State
  - PC Program Counter
  - Register File
    - Heavily used data
  - Condition Codes

## □ Memory

- Byte array
- Code + data
- stack

# Why Different Processors are Used?

- What is the difference between processors used in desktops, laptops, mobile phones, washing machines etc. ?
  - Performance/speed
  - Power consumption
  - Cost
  - General purpose/special purpose
  - Domain Specific: Network, DSP, Image.Crypto

# The Embedded Processor

## What is it?

A programmable processor whose programming interface is not accessible to the end-user of the product.

The only user-interaction is through the actual application.


## Examples:

- Sharp PDA's are encapsulated products with fixed functionality
- 3COM Palm pilots were originally intended as embedded systems. Opening up the programmers interface turned them into more generic computer systems.

# Some interesting numbers

- The Intel 4004 was intended for an embedded application (a calculator)
- Of today's microprocessors
  - 95% go into embedded applications
    - SSH3/4 (Hitachi): best selling RISC microprocessor
  - 50% of microprocessor revenue stems from embedded systems
- Often focused on particular application area
  - Microcontrollers
  - DSPs
  - Media Processors
  - Graphics Processors
  - Network and Communication Processors



# Architecture Design: Measurement and Evaluation



Creativity

Architecture Design is an iterative process:

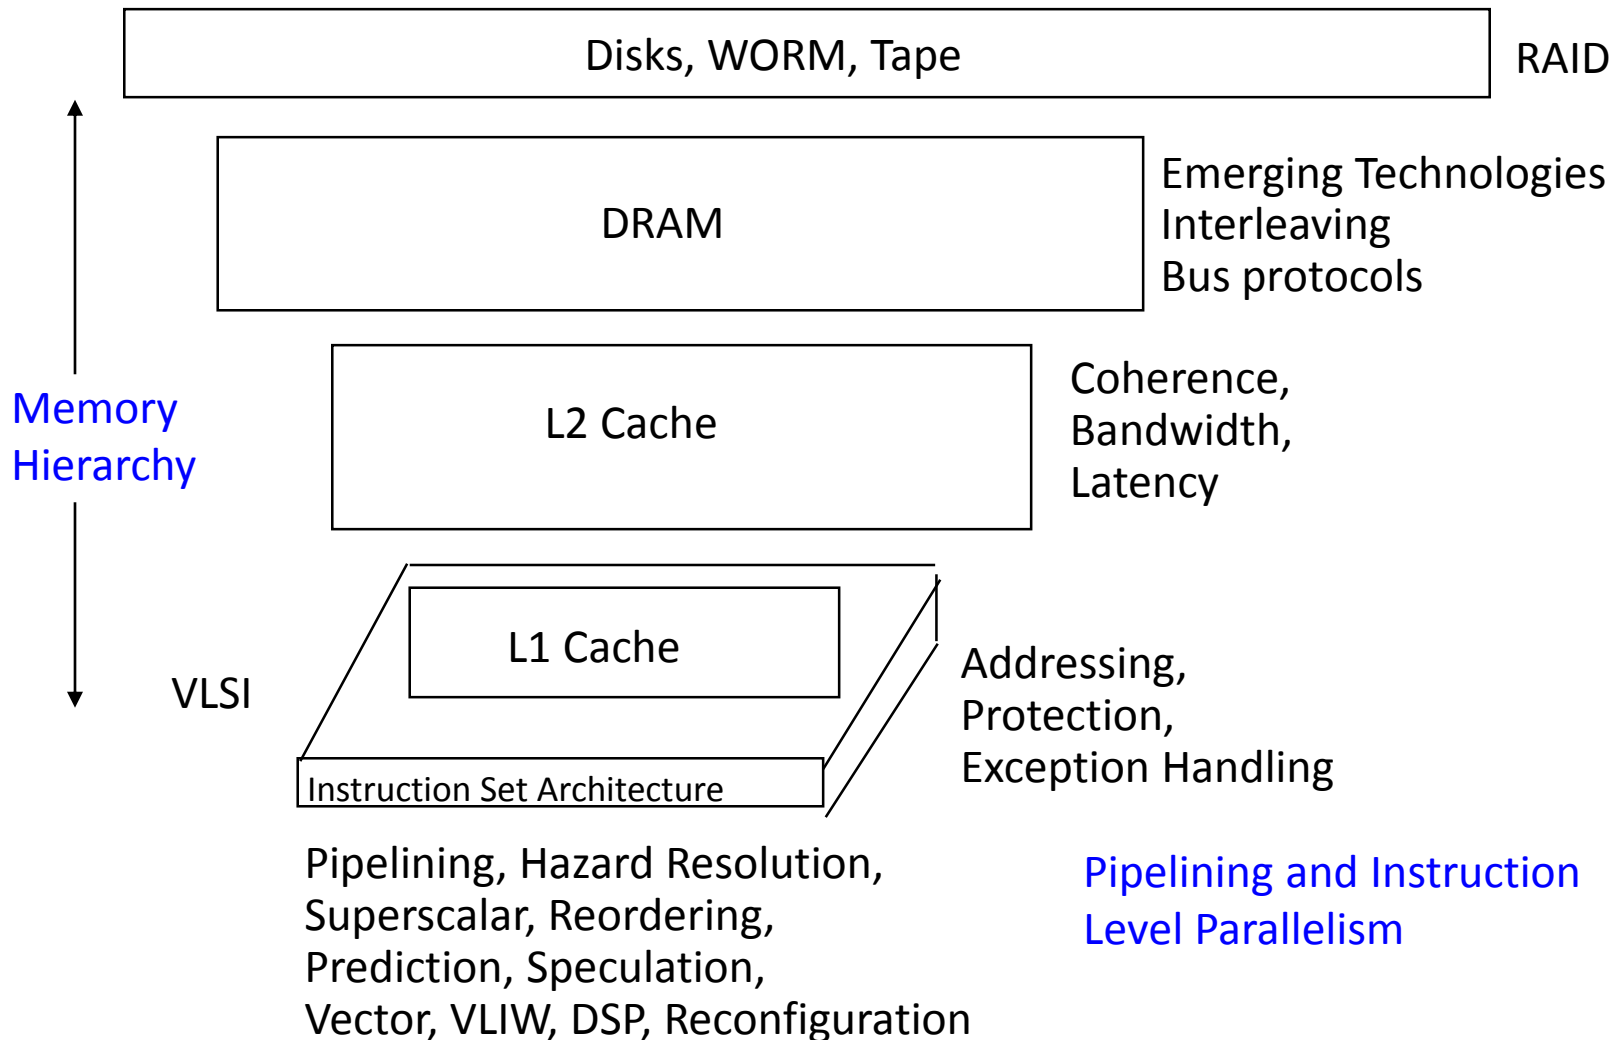
- Searching the space of possible designs
- At all levels of computer systems




Bad Ideas

# Computer Architecture Topics

## Input/Output and Storage



# Computer Architecture Topics



Shared Memory,  
Message Passing,  
Data Parallelism

Network Interfaces


Processor-Memory-Switch

Multiprocessors

Networks and Interconnections

Topologies,  
Routing,  
Bandwidth,  
Latency,  
Reliability

# Computer Engineering Methodology






# Measurement Tools

- Hardware: Cost, delay, area, power estimation
- Benchmarks, Traces, Mixes
- Simulation (many levels)
  - ISA, RT, Gate, Circuit
- Queuing Theory
- Rules of Thumb
- Fundamental “Laws”/Principles

# Metric 1: Performance

In passenger-mile/hour



Plane	DC to Paris	Speed	Passengers	Throughput
Boeing 747	6.5 hours	610 mph	470	286,700
Concorde	3 hours	1350 mph	132	178,200

- Time to run the task
  - Execution time, response time, latency
- Tasks per day, hour, week, sec, ns ...
  - Throughput, bandwidth

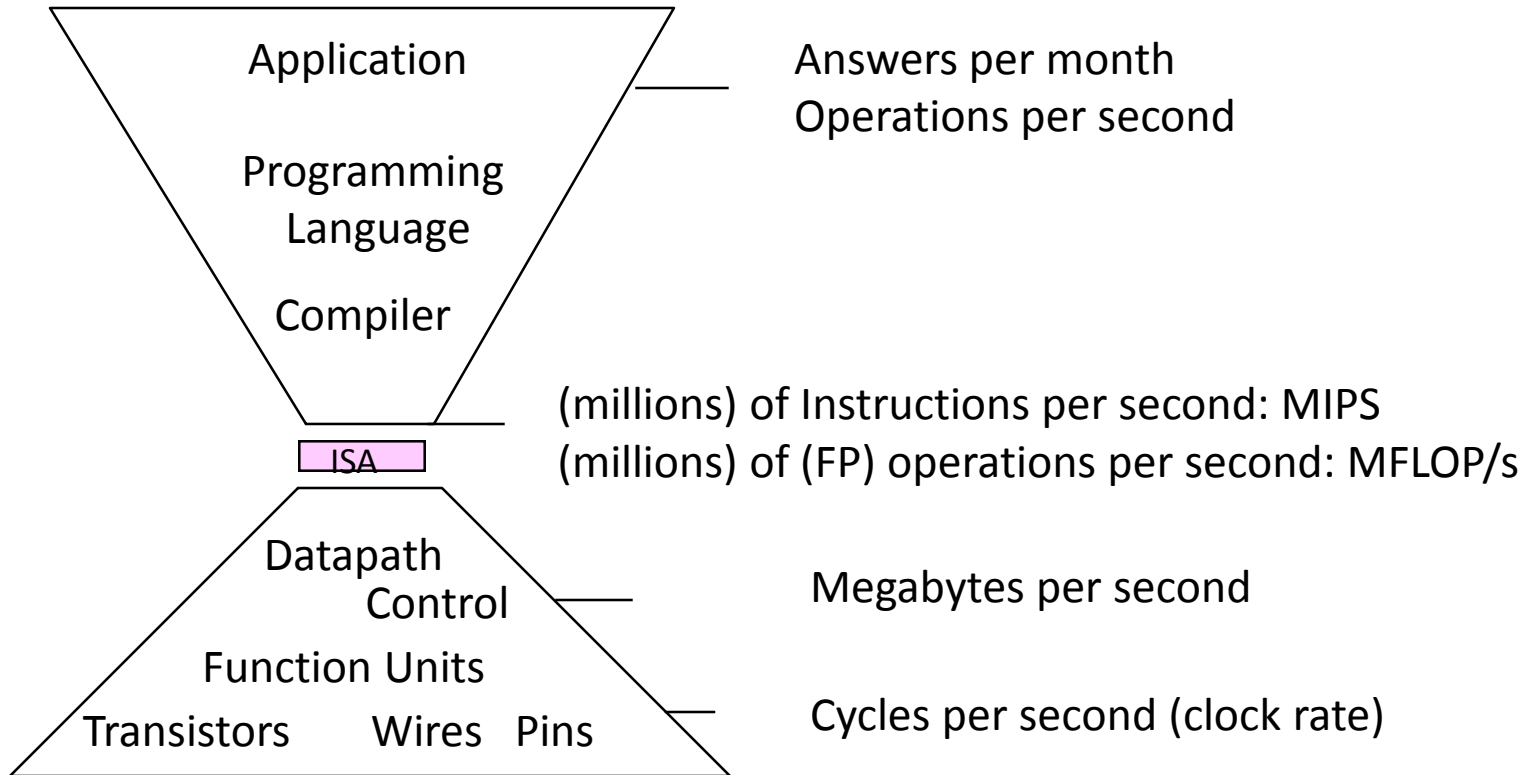
# The Performance Metric

"X is n times faster than Y" means

$$\frac{\text{ExTime}(Y)}{\text{ExTime}(X)} = \frac{\text{Performance}(X)}{\text{Performance}(Y)}$$

- Speed of Concorde vs. Boeing 747
- Throughput of Boeing 747 vs. Concorde

# Metrics of Performance



# Creating Benchmark Sets

- Real programs
- Kernels
- Toy benchmarks
- Synthetic benchmarks
  - e.g. Whetstones and Dhrystones

# SPEC: System Performance Evaluation

## Cooperative

- First Round 1989
  - 10 programs yielding a single number (“SPECmarks”)
- Second Round 1992
  - SPECInt92 (6 integer programs) and SPECfp92 (14 floating point programs)
    - Compiler Flags unlimited. March 93 of DEC 4000 Model 610:  
spice: `unix.c:/def=(sysv,has_bcopy,"bcopy(a,b,c)=memcpy(b,a,c) "`  
wave5: `/ali=(all,dcom=nat)/ag=a/ur=4/ur=200`  
nasa7: `/norecu/ag=a/ur=4/ur2=200/lc=blas`
- Third Round 1995
  - new set of programs: SPECInt95 (8 integer programs) and SPECfp95 (10 floating point)
  - “benchmarks useful for 3 years”
  - Single flag setting for all programs: SPECInt\_base95, SPECfp\_base95

# Performance Evaluation


- “For better or worse, benchmarks shape a field”
- Good products created when have:
  - Good benchmarks
  - Good ways to summarize performance
- Given sales is a function in part of performance relative to competition, investment in improving product as reported by performance summary
- If benchmarks/summary inadequate, then choose between improving product for real programs vs. improving product to get more sales;  
Sales almost always wins!
- Execution time is the measure of computer performance!

# Integrated Circuits Costs

$$\text{IC cost} = \frac{\text{Die cost} + \text{Testing cost} + \text{Packaging cost}}{\text{Final test yield}}$$

$$\text{Die cost} = \frac{\text{Wafer cost}}{\text{Dies per Wafer} \times \text{Die yield}}$$

$$\text{Dies per wafer} = \frac{\pi (\text{Wafer\_diam}/2)^2}{\text{Die\_Area}} - \frac{\pi \times \text{Wafer\_diam}}{\sqrt{2} \cdot \text{Die\_Area}} - \text{Test\_Die}$$




$$\text{Die Yield} = \text{Wafer\_yield} \times \left\{ 1 + \left( \frac{\text{Defect\_Density} \times \text{Die\_area}}{\alpha} \right)^{-\alpha} \right\}$$

Die Cost goes roughly with die area



# Power/Energy



- Lead processor power increases every generation
- ✕ Compactions provide higher performance at lower power